



Computer Networks

Wenzhong Li

Nanjing University

Fall 2014



Chapter 5. End-to-End Protocols

- Transport Services and Mechanisms
- User Datagram Protocol (UDP)
- Transmission Control Protocol (TCP)
- TCP Congestion Control
- Real-time Transport Protocol (RTP)
- Session Initiation Protocol (SIP)
- Real Time Streaming Protocol (RTSP)



User Datagram Protocol (UDP)

- User datagram protocol, RFC 768
- **Connectionless service** for application level processes
 - Unreliable, “best-effort” of IP
 - Each UDP segment handled independently of others
 - Delivery and duplication control not guaranteed
- **Simple and reduced overhead**
 - No connection establishment
 - No connection state at sender, receiver
 - Small segment header

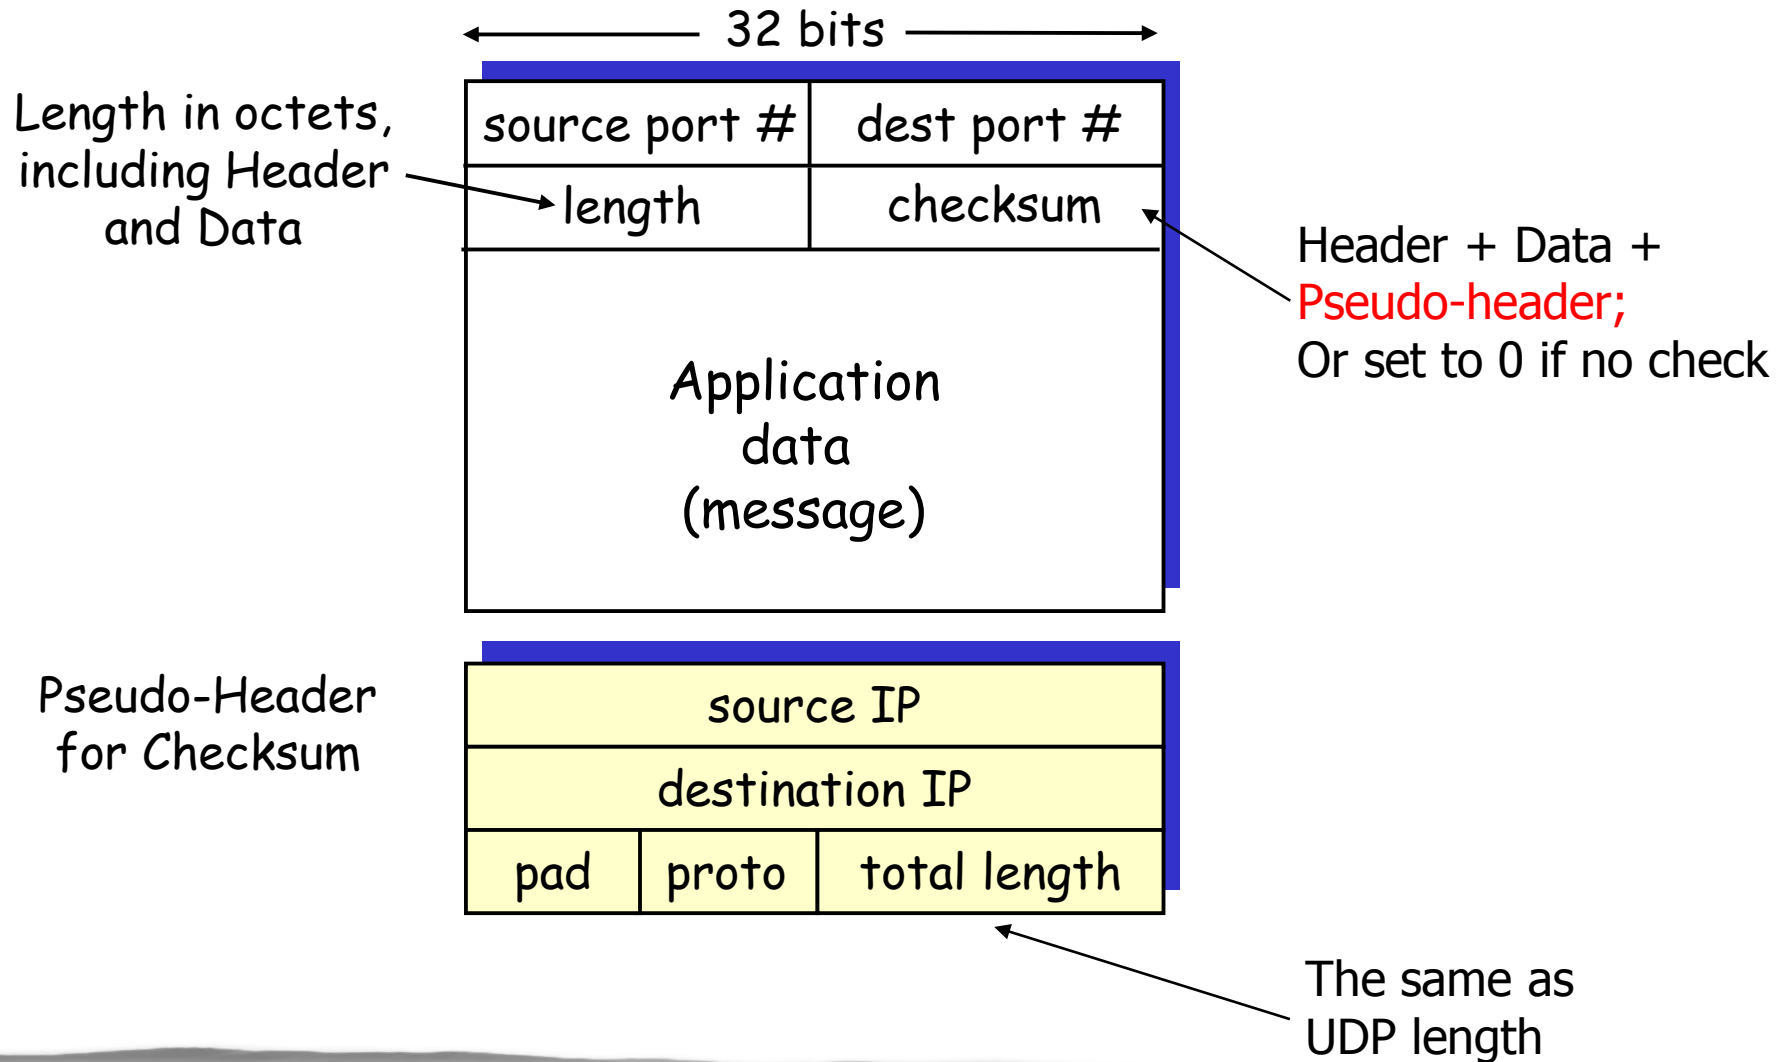


UDP Uses

- Normal use
 - Inward data collection from sensors
 - Outward data dissemination
 - Real time applications
 - Request-Response (e.g. RPC), add reliability at application layer
- Example Apps based on UDP
 - DNS
 - SNMP



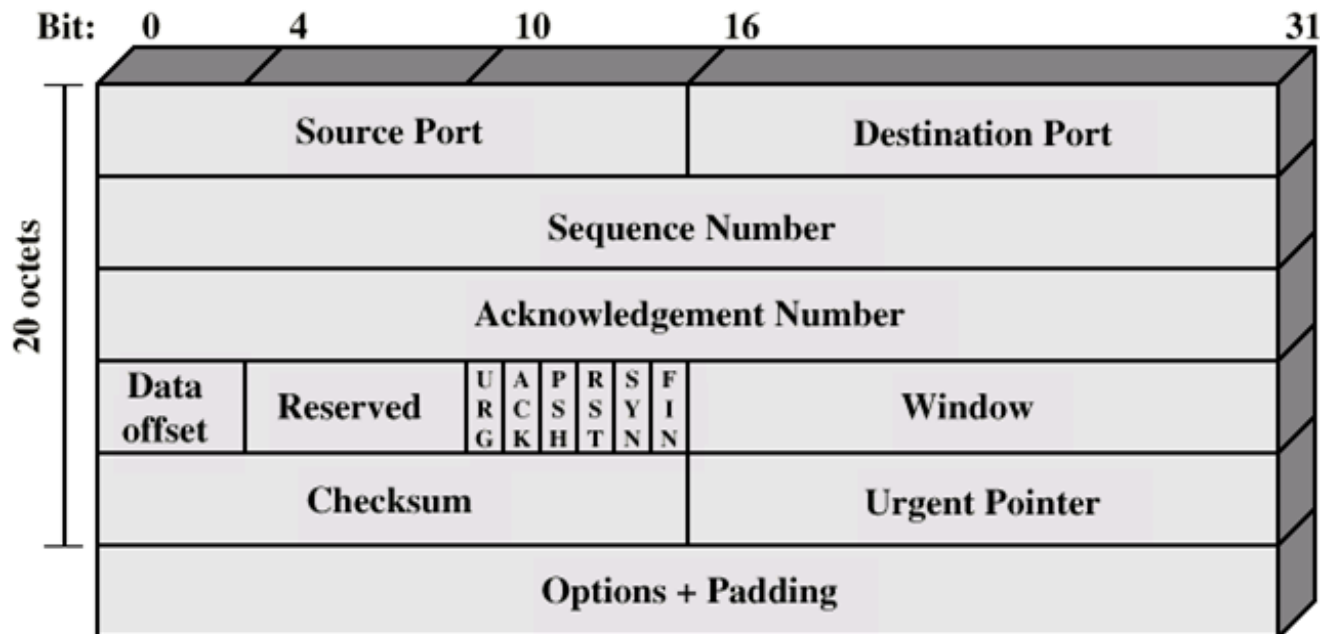
UDP Segment Format





Transmission Control Protocol (TCP)

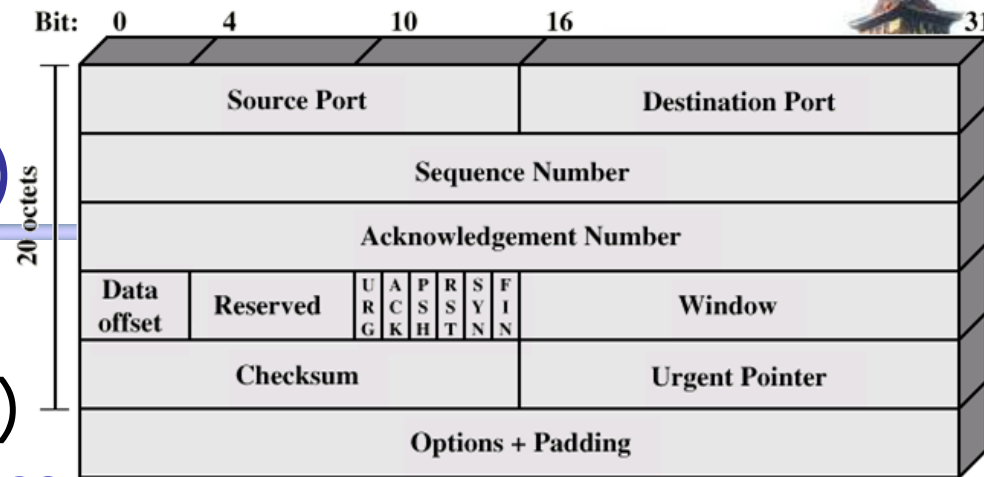
- **Reliable communication** between pairs of processes
 - Across variety of reliable and unreliable networks and internets
 - 20 bytes





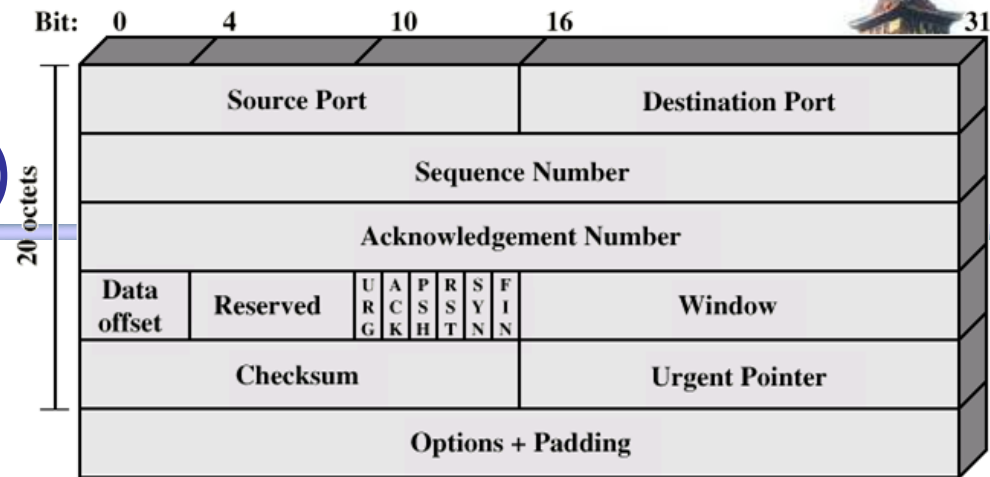
TCP Header Fields (1)

- Source port (16 bits)
- Destination port (16 bits)
 - Identify src and dest TCP user
- Sequence number (32 bits)
 - Seq number of first data octet
 - If SYN is set, it is ISN and first data octet is ISN+1
- ACK number (32 bits)
 - Piggybacked ACK
- Window (16 bits)
 - Credit allocation in octets, i.e. rcv_window of sender





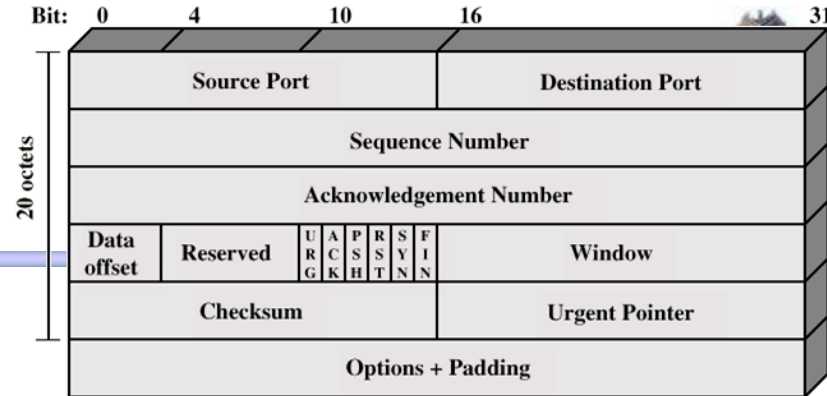
TCP Header Fields (2)



- Data offset (4 bits)
 - Number of 32-bit words in the header (报头长度)
 - Largest data offset is $15 \times 4 = 60$ octets
- Checksum (16 bits)
 - Header + Data + Pseudo-header (src IP, dest IP, protocol No, total length)
- Reserved (6 bits)
- Options (Variable)
 - e.g. Maximum segment size the sender can accept
 - or Max value of `rcv_window`



TCP Header Fields (3)



■ Flags (6 bits):

- URG: urgent pointer field meaningful
- ACK: acknowledgment field meaningful
- PSH: push function
- RST: reset the connection
- SYN: synchronize the sequence number
- FIN: no more data from sender

■ Urgent Pointer (16 bits)

- Points to last octet in a sequence of urgent data



Parameters Passed to IP

- TCP passes **QOS parameters** down to IP
 - Precedence
 - Normal delay / low delay
 - Normal throughput / high throughput
 - Normal reliability / high reliability
- IPv4 "Type of Service" or IPv6 "Traffic Class"



TCP Service Request Primitives

Primitive	Parameters	Description
Unspecified Passive Open	source-port, [timeout], [timeout-action], [precedence], [security-range]	Listen for connection attempt at specified security and precedence from any remote destination.
Fully Specified Passive Open	source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security-range]	Listen for connection attempt at specified security and precedence from specified destination.
Active Open	source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security]	Request connection at a particular security and precedence to a specified destination.
Active Open with Data	source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security], data, data-length, PUSH-flag, URGENT-flag	Request connection at a particular security and precedence to a specified destination and transmit data with the request
Send	local-connection-name, data, data-length, PUSH-flag, URGENT-flag, [timeout], [timeout-action]	Transfer data across named connection
Allocate	local-connection-name, data-length	Issue incremental allocation for receive data to TCP
Close	local-connection-name	Close connection gracefully
Abort	local-connection-name	Close connection abruptly
Status	local-connection-name	Query connection status



TCP Service Response Primitives

Primitive	Parameters	Description
Open ID	local-connection-name, source-port, destination-port*, destination-address*,	Informs TCP user of connection name assigned to pending connection requested in an Open primitive
Open Failure	local-connection-name	Reports failure of an Active Open request
Open Success	local-connection-name	Reports completion of pending Open request
Deliver	local-connection-name, data, data-length, URGENT-flag	Reports arrival of data
Closing	local-connection-name	Reports that remote TCP user has issued a Close and that all data sent by remote user has been delivered
Terminate	local-connection-name, description	Reports that the connection has been terminated; a description of the reason for termination is provided
Status Response	local-connection-name, source-port, source-address, destination-port, destination-address, connection-state, receive-window, send-window, amount-awaiting-ACK, amount-awaiting-receipt, urgent-state, precedence, security, timeout	Reports current status of connection
Error	local-connection-name, description	Reports service-request or internal error

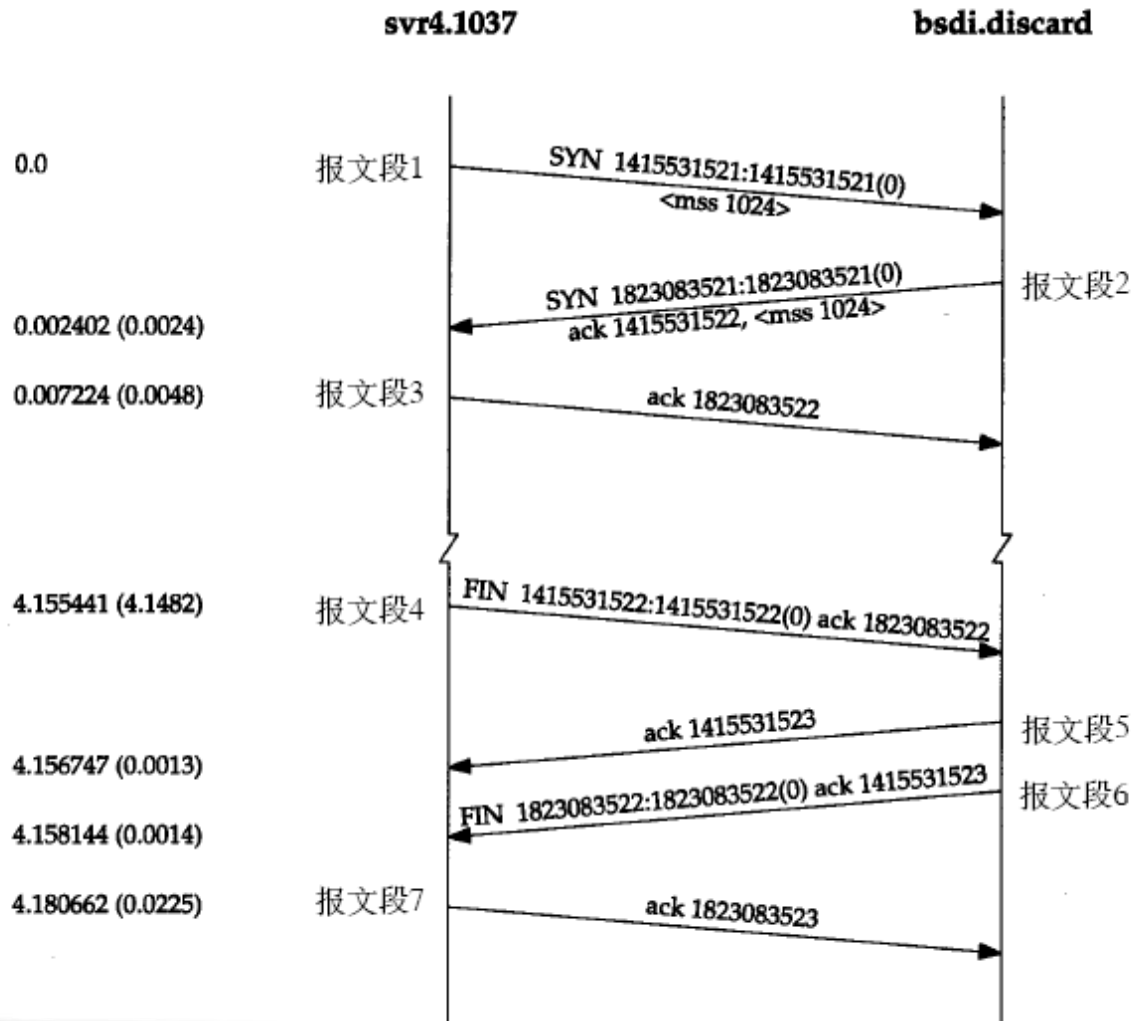


TCP Mechanisms (1)

- Connection establishment
 - 3-way handshake
 - Between pairs of ports
 - One port can connect to multiple destination ports
- Connection termination
 - Graceful termination: `CLOSE + FIN`
 - Abrupt termination: `ABORT + RST`



Illustration of Connection establishment & Termination





TCP Mechanisms (2)

■ Data transfer

- **Logical stream** of octets
- Octets numbered modulo 2^{32}
- Flow control by credit allocation of number of octets
- **Data buffered** at sender and receiver
- User sets **PUSH** to force data transmission immediately
- User may specify a block of data as **urgent**



Implementation Policy

- Send
- Deliver
- Accept
- Retransmit
- Acknowledge



Send

- If no *PUSH* or *CLOSE*, TCP entity transmits at its own convenience
- Data issued by TCP user buffered at transmit buffer
 - May construct segment per data batch
 - May wait for certain amount of data



Deliver

- In absence of *PUSH*, TCP entity delivers data at own convenience
- May deliver as each segment in order received
 - Deliveries (I/O interrupts) are frequent and small
- May buffer data from more than one segment
 - Deliveries are infrequent and large



Accept

- Segments may arrive **out of order**
- In order
 - Only accept segments in order
 - Discard out of order segments
 - Makes for a simpler implementation
- In windows
 - Accept all segments **within receive window**
 - Can reduce retransmission



Retransmit

- TCP entity maintains queue of segments **transmitted but not acknowledged**
- TCP will retransmit if not ACKed in given time
 - **First only**: one timer a queue, reset the timer after retransmission of first segment in queue
 - **Batch**: one timer a queue, reset after retransmission of all segments in queue
 - **Individual**: one timer each segment, reset after retransmission



Fast Retransmit

- Time-out period often relatively long
- Detect lost segments via **duplicate ACKs**
 - If segment is lost, there will likely be many duplicate ACKs
- If a TCP entity receives 3 ACKs for the same data, then segments after ACKed data must be lost
 - Trigger **fast retransmit**: resend segment before timer expires



Acknowledgement

■ Immediate or Cumulative

Event at Receiver

TCP Receiver action

Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed

Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK

Arrival of in-order segment with expected seq #. One other segment has ACK pending

Immediately send single cumulative ACK, ACKing both in-order segments

Arrival of out-of-order segment higher-than-expect seq. # . Gap detected

Immediately send duplicate ACK, indicating seq. # of next expected byte

Arrival of segment that partially or completely fills gap

Immediate send ACK, provided that segment sat at lower end of gap



TCP congestion control



TCP Congestion Control

- Congestion control
 - Too many sources sending too much data too fast for Internet to handle
 - RFC 1122, Requirements for Internet hosts
- End to end control, **no network assistance**
 - Retransmission timer
 - Window management



Retransmission Timer Management

- Estimate round trip delay by **observing delay pattern**
 - Simple average
 - Exponential average
- **Set timer** to value somewhat greater than estimate
 - RFC 793
 - RTT Variance Estimation (Jacobson's algorithm)
- How to **set timer after retransmission**
 - Exponential RTO backoff algorithm
- **When to sample** the round trip delay
 - Karn's Algorithm



Simple Average

■ Term

- $RTT(i)$: round-trip time observed for the i^{th} transmitted segment
- $ARTT(k)$: **average round-trip time** for the first k segments

■ Expression

$$ARTT(k+1) = \frac{1}{k+1} \sum_{i=1}^{k+1} RTT(i) \quad \text{or}$$

$$ARTT(k+1) = \frac{1}{k+1} (k \times ARTT(k) + RTT(i))$$

$$\frac{k}{k+1} ARTT(k) + \frac{1}{k+1} RTT(i)$$




Exponential Average

- Term
 - $SRTT(k)$: **smoothed round-trip time** estimate for the first k segments
- Expression


$$SRTT(k+1) = \alpha \times SRTT(k) + (1-\alpha) \times RTT(k+1) \quad \text{i.e.}$$

$$SRTT(k+1) = (1-\alpha) \times RTT(k+1) + \alpha(1-\alpha) \times RTT(k) + \\ \alpha^2(1-\alpha) \times RTT(k-1) + \dots + \alpha^k(1-\alpha) \times RTT(1)$$

Simple and Exponential Averaging



(a) Increasing function



(b) Decreasing function



RFC 793

■ Term

- $RTO(k)$: **retransmission timeout**, i.e. the timer after the first k segments

■ Expression

$$RTO(k+1) = \text{Min}(\text{UBOUND}, \text{MAX}(\text{LBOUND}, \beta \times SRTT(k+1)))$$

- Retransmission timer set between $\text{LBOUND} \sim \text{UBOUND}$
- Suggested values, α : 0.8~0.9, β : 1.3~2.0



Jacobson's Algorithm (1)

Problem in RFC 793

- Not counting **variance of RTT** (network stability)
- When network is stable, RTT variance is low, but $\beta=1.3$ gives a higher RTO
- When network is **unstable**, RTT variance is high, $\beta=2$ is inadequate to protect against retransmissions



Jacobson's Algorithm (2)

■ Term

- $SERR(k)$: **smoothed error estimate**, difference of round-trip time of segment k and the current $SRTT$
- $SDEV(k)$: **standard deviation** for round-trip time of first k segments

■ Expression

$$SRTT(k+1) = (1-g) \times SRTT(k) + g \times RTT(k+1)$$


$$SERR(k+1) = RTT(k+1) - SRTT(k)$$

$$SDEV(k+1) = (1-h) \times SDEV(k) + h \times |SERR(k+1)|$$


$$RTO(k+1) = SRTT(k+1) + f \times SDEV(k+1)$$

$$g = 1/8 = 0.125 \quad h = 1/4 = 0.25 \quad f = 2 \text{ or } 4$$

Jacobson's RTO Calculation



(a) Increasing function



(b) Decreasing function



Exponential RTO Backoff

- Timeout is often **due to congestion** by dropped packet or long round trip
 - Should slow down end system transmission
 - Maintaining RTO is not a good idea
- Similar to **Binary exponential backoff** in Ethernet
 - RTO multiplied each time a segment is re-transmitted
 - $RTO = q \times RTO$
 - Commonly $q=2$



Karn's Algorithm

- The problem
 - If a segment is **re-transmitted**, the *ACK* arriving may be
 - For the first copy of the segment, or for the second copy, or others
 - No way to tell
- RTT Sampling
 - Do **not measure RTT** for re-transmitted segments
 - Calculate RTO backoff when re-transmission occurs
 - Until *ACK* arrives for segment that has not been re-transmitted
 - Begin sampling, stop RTO backoff




Window Management

- Add **congestion (send) window** besides the credit (receive window)
 $awind = Min(credit, cwnd)$
 - **awnd**: allowed window, in MSS (maximum segment size)
 - **credit**: the amount of unused credit granted in last ACK, in MSS
 - **cwnd**: congestion window, in MSS
- Manage congestion window
 - **Slow Start**: exponentially expending the cwnd at start of connection
 - **Dynamic Window on Congestion**: shrinking / expending the cwnd with stages when retransmission occurs



Slow Start

- When connection begins, **cwnd = 1 MSS**
- Each time an *ACK* received, cwnd increased by ACKed number of MSSs until Max value reached
- cwnd increased exponentially until first loss event occurs
 - Timeout or 3 duplicate *ACK*s





Dynamic Windows Sizing

- By Jacobson, set slow start threshold **ssthresh** = $cwnd/2$
- After **3 duplicate ACKs**
 - Network still capable of delivering some segments
 - $cwnd$ is set to be **ssthresh**
 - $cwnd$ increases by 1 (linearly instead of exponentially) after each RTT or **ACK** received
- After **timeout event**
 - $cwnd$ is set to 1
 - Same as **Slow Start** until $cwnd$ reaches **ssthresh**
 - $cwnd$ increases by 1 after each RTT or **ACK** received



Illustration of Window Management




Figure 17.14 Illustration of Slow Start and Congestion Avoidance



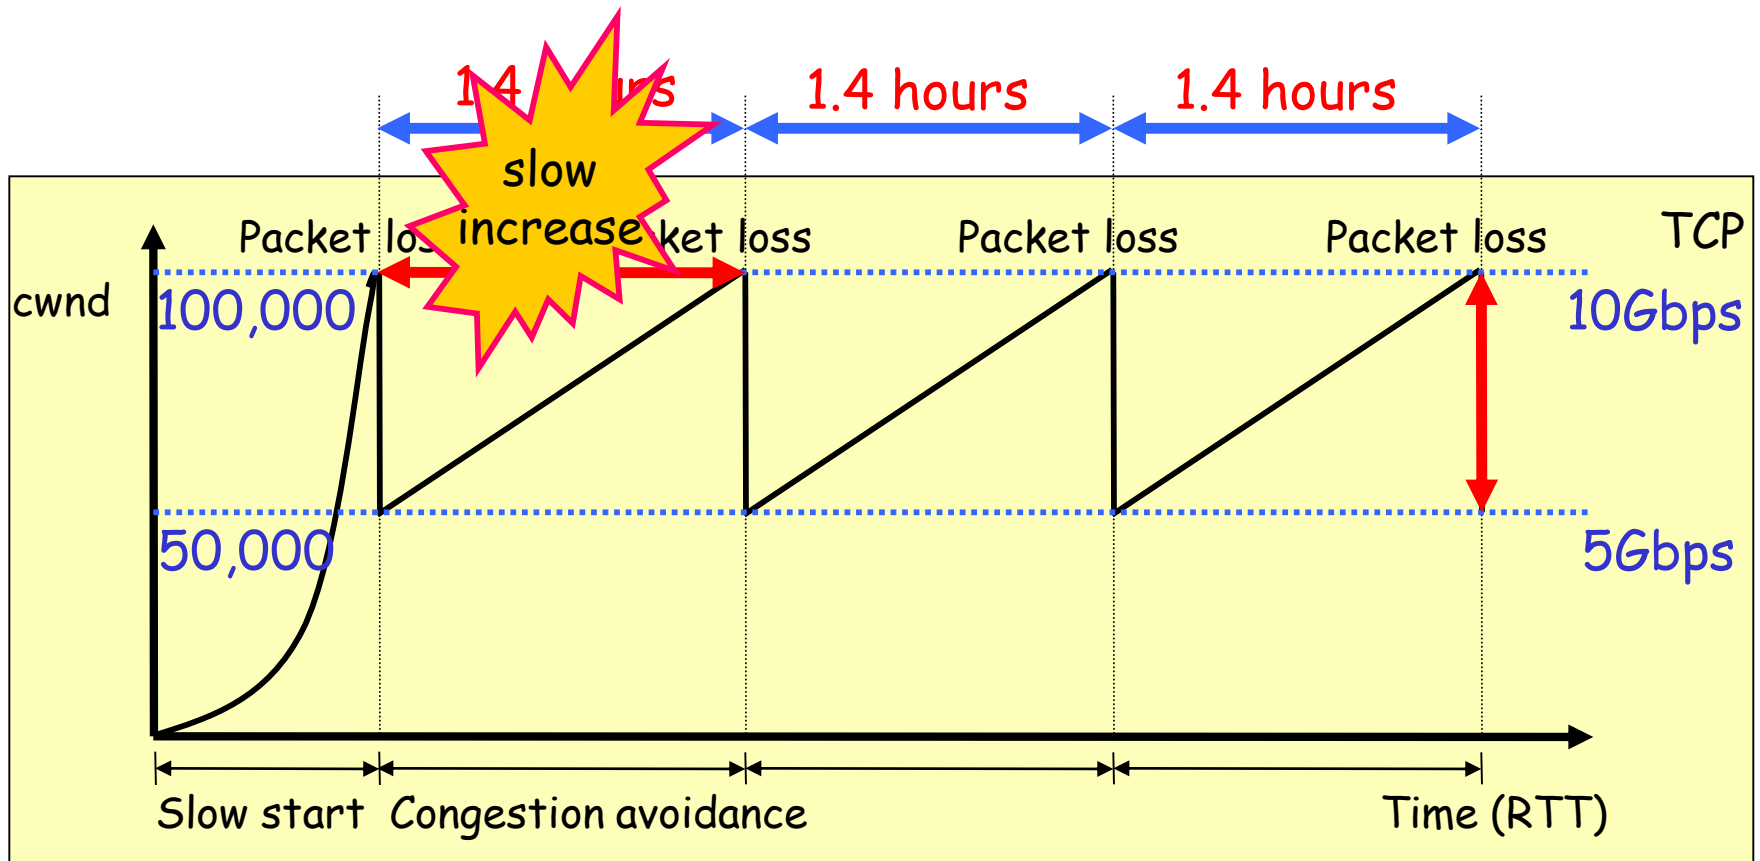
New Window Management Algorithm

- Problem of **linear increase**
 - Many **long fat networks**: large bandwidth with long delay
 - Slow response of TCP in such networks leaves sizeable unused bandwidth
- An example
 - A TCP connection with *1250-Octet MSS* (*Maxitum Segment Size*) and *100ms RTT* on *10Gbps* network
 - To fully use the network, credit is big, and nearly **1.4 hour** is needed for linear increase



Too Slow Linear Increase

- Q: how to improve it





BIC and CUBIC

- **BIC** (Binary Increase Congestion control)
 - Implemented and used by default in Linux kernels 2.6.8
- **CUBIC**
 - The window is a cubic function of time since the last congestion event
 - Implemented and used by default in Linux kernels 2.6.19 and above



- BIC adaptively increase cwnd, and decrease cwnd by 1/8

$$cwnd = cwnd + 1$$




$$cwnd = cwnd + f(cwnd, history)$$

$$cwnd = cwnd * (1 - 1/2)$$



$$cwnd = cwnd * (1 - 1/8)$$





BIC Overview

- 2 stages
 - **Binary Search**: increase window after congestion
 - **Max Probing**: search for better window size (until max credit)
 - (Max window size if set from history information. If there are more size available, Max_Probing is used for the purpose of exploring larger window size)
- 4 parameters defined
 - ***S_{max}*** : the maximum increment, e.g. $1/8 \times credit$
 - ***S_{min}*** : the minimum increment, e.g. 2 MSS
 - (Using S_{max} , S_{min} to avoid jitter: If a binary search steps too large, the traffic will change quickly. Thus S_{max} is used to limit the maximum change in one step)
 - ***W_{max}*** : maximum window size of current search, e.g. window size just before the lost
 - ***W_{min}*** : minimum window size of current search, e.g. current window size without lost



BIC Stages

■ Additive Increase

- Linear increase with inc


■ Binary search

- How to set inc

```
while ( $Wmin \leq Wmax$ ){  
     $inc = (Wmin + Wmax) / 2 - cwnd$ ;  
    if ( $inc > Smax$ )  $inc = Smax$ ;  
    else if ( $inc < Smin$ )  $inc = Smin$ ;  
     $cwnd += inc$ ;  
  
    if (no packet losses)  $Wmin = cwnd$ ;  
    else {  
         $Wmax = cwnd$ ;  
         $Wmin = cwnd \times \beta$  (e.g. 0.8)  
    }  
}
```

BIC Stages

- **Binary Search Stage:** Additive increase + Binary search
- **Max Probing Stage:** Binary search + Additive increase



At the beginning, each step increases S_{\max}

In the middle, pure binary search

To probing more window size, each step increases by binary search or S_{\min}



CUBIC

- BIC problem
 - The BIC's growth function may be **too aggressive** for TCP
 - BIC is not suitable for short RTT or low speed networks (may cause unfairness)
- Handle
 - Express the multi-stage BIC curve with a **single cubic function**



CUBIC Overview

■ Parameters


- ***W_{cubic}***: current cwnd
- ***W_{max}***: window size just before the last lost
- ***T***: elapsed time from the last lost
- ***C***: a scaling constant

■ Function $W_{cubic} = C(T - K)^3 + W_{max}$

$$\text{where } K = \sqrt[3]{\frac{W_{max} \times (1 - \beta)}{C}}$$



CUBIC Overview





Summary

- UDP & TCP
- TCP header fields
- TCP congestion control
 - Retransmission timer
 - Window management



Homework

- 书第22章习题: 22.3, 22.4, 22.9, 22.14, 22.16, 22.17, 22.18, 22.21